# Found in the Middle: How Language Models Use Long Contexts Better via Plug-and-Play Positional Encoding

**Zhenyu Zhang** [1][♣]  **Runjin Chen** [1]  **Shiwei Liu** [2]  **Zhewei Yao** [3]  **Olatunji Ruwase** [3]  **Beidi Chen** [4]  **Xiaoxia Wu** [3][♠]  **Zhangyang Wang** [1][♠]

## Abstract

This paper aims to overcome the "lost-in-the-middle" challenge of large language models (LLMs). While recent advancements have successfully enabled LLMs to perform stable language modeling with up to 4 million tokens, the persistent difficulty faced by most LLMs in identifying relevant information situated in the middle of the context has not been adequately tackled. To address this problem, this paper introduces Multi-scale Positional Encoding (Ms-PoE) which is a simple yet effective plug-and-play approach to enhance the capacity of LLMs to handle the relevant information located in the middle of the context, without fine-tuning or introducing any additional overhead. Ms-PoE leverages the position indice rescaling to relieve the long-term decay effect introduced by RoPE, while meticulously assigning distinct scaling ratios to different attention heads to preserve essential knowledge learned during the pre-training step, forming a multi-scale context fusion from short to long distance. Extensive experiments with a wide range of LLMs demonstrate the efficacy of our approach. Notably, Ms-PoE achieves an average accuracy gain of up to 3.8 on the Zero-SCROLLS benchmark over the original LLMs. Code are available at `https://github.com/VITA-Group/Ms-PoE`

## 1. Introduction

Effective long-sequence reasoning in large language models (LLMs) is crucial for a wide range of applications (Ré et al., 2022; Li et al., 2023), from understanding extensive texts (Tay et al., 2020; Kryściński et al., 2021) and manag-

♣ Work done during internship at Microsoft. [1]The University of Texas at Austin [2]University of Oxford [3]Microsoft [4]Carnegie Mellon University. ♠ Correspondence to: Xiaoxia Wu <xiaoxiawu@microsoft.com>, Zhangyang Wang <atlaswang@utexas.edu>.
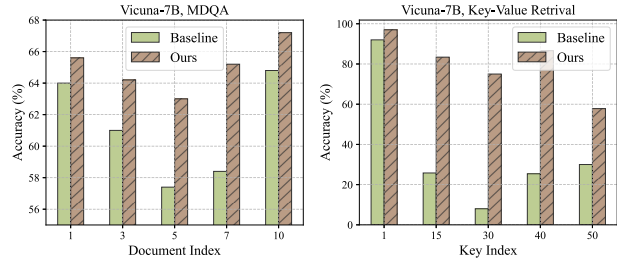
*Preprint.*



*Figure 1.* The x-axis illustrates the placement of essential information within the prompt, ranging from start to end. The green bar serves as a standard baseline, illustrating the "lost-in-the-middle" phenomenon. We introduce our method, **M**ulti-**s**cale **Po**sition **E**ncoding (Ms-PoE), which requires neither additional fine-tuning nor increased memory usage. Instead, it involves a simple remapping of the position embedding depicted in Figure 2, which enables the important information in the middle to be detected effectively (brown bars). For more details, see Section 4.2 and Figure 5.

ing day-long conversations (Zhang et al., 2021; Zhong et al., 2022) to code generation (Du et al., 2023; Zheng et al., 2023) and science discoveries (Varadi et al., 2022; Song et al., 2023b). Recent system support advancements (Dao, 2023; Jacobs et al., 2023) have enabled training transformers for any $L$ sequence length even with $O(L^2)$ computational complexity. This is exemplified by models such as MPT (Team, 2023) and Mistral (Jiang et al., 2024) pre-trained with sequence lengths 16k and 32k respectively.

Nevertheless, emerging research reveals the constrained efficacy of LLMs in managing tasks requiring long contextual understanding. Particularly, Liu et al. (2023) demonstrated a substantial degradation in LLMs' performance when crucial information is positioned amidst a lengthy context, a phenomenon they refer to as "lost-in-the-middle". One explanation is about the use of rotary positional embedding (RoPE) (Su et al., 2024), a prevalent positional encoding technique used in open-source LLMs. As a relative position embedding, RoPE incorporates a long-term decay property, predisposing the model to prioritize current/nearby tokens while paying less attention to further ones. Xiao et al. (2023) identified a surprising trend attributed to the Softmax operation where attention scores are disproportionately allocated

into initial tokens, irrespective of their relevance to the language modeling task. Despite the presence of considerable redundancy in long-context inputs (Zhang et al., 2023), crucial information may be located across different positions. The inclination of LLMs to overlook the middle section presents a challenge for their applications, particularly in the context of long-context reasoning. Several approaches successfully extend pre-trained LLMs with context up to extreme token length, either through sparse selection of crucial tokens during generation (Xiao et al., 2023; Zhang et al., 2023; Han et al., 2023) or by modifying positional encoding (Chen et al., 2023c; Jin et al., 2024). *Nevertheless, these approaches primarily aim to extend the context length of LLMs and, consequently, fall short in addressing the "lost-in-the-middle" problem when applied out-of-the-box.*

Efforts have been made to enhance LLMs' capacity to capture vital information located within the middle of the context. These include extra memory bank (Wang et al., 2023), reordering the input context based on relevance (Peysakhovich & Lerer, 2023; Chen et al., 2023d), enhancing the information searching and reflection ability via attention strengthening tasks (Junqing et al., 2023; Xu et al., 2023), splitting the input into short segments and applying short-text models (Ivgi et al., 2023). For example, Peysakhovich & Lerer (2023) empirically discovered that LLMs tend to emphasize more on the current window while still paying more attention to the relevant text than distracting content. They subsequently introduced "attention sorting" where the main idea is iteratively sorting documents based on their attention scores, such that critical information will likely be placed at the end, to fit the position-biased nature of RoPE. Chen et al. (2023d) conducted parallel runs of LLMs with different RoPE angles, thereby mitigating the risk of overlooking crucial information through a weighted sum of the outputs. These approaches usually require additional memory or multiple inference runs, which can be expensive for LLMs.

In this paper, we aim to address the "lost-in-the-middle" problem by reintroducing the concept of multi-scale features from computer vision into the context of Transformer-based LLMs. Multi-scale features, well-established in Inception-style models (Szegedy et al., 2015; 2016; Guo et al., 2022), utilize parallel employment of kernels with different sizes to fuse multi-scale information, spanning short to long distances. Introducing multi-scale operations into LLMs intuitively can help compensate for crucial information located in the middle, which might be easily overlooked by full attention operation. Unlike modifying the attention module to form multi-scale attention, we choose to re-scale the indices of positional encoding. This decision is grounded not only in its effectiveness in easily adjusting the scale of the context window by simply changing the position indices (Chen et al., 2023c) but also in the potential of down-scaling

the position indices to relieve the long-term decay property introduced by RoPE. However, this approach was initially introduced to extend context windows, and its performance regarding the "lost-in-the-middle" problem remains uncertain for several reasons: (i) Indice re-scaling forces position embeddings of original context window to reside in a narrower region, leading to performance degradation in the original context window as shown in Chen et al. (2023c). (ii) Uniformly applying the same scaling ratio throughout the entire model might be sub-optimal to preserve essential knowledge learned during pre-training; (ii) Fine-tuning is necessary for the original approach, albeit minimal. The impact without fine-tuning remains unknown.

To this end, we systematically visit the position indices scaling regarding the "lost-in-the-middle" problem and counter-intuitively discover that it is possible to slightly mitigate the "lost-in-the-middle" issue if we carefully choose the scaling ratio to be around 1.5-2. Additionally, we observe that different attention heads exhibit varying sensitivity to the position shift of the relevant document. Some attention heads are "position-aware", consistently capturing relevant information even with position shifts, while others may occasionally capture position changes, and some heads are completely insensitive to position changes. This highlights the need to treat attention heads separately when re-scaling position indices.

**Contribution.** Inspired by the above observations, we introduce Multi-scale Positional Encoding (Ms-PoE), a simple yet effective plug-and-play approach that can enhance the long-context reasoning capability of pre-trained LLMs without requiring fine-tuning or introducing any additional overhead. Ms-PoE meticulously assigns distinct scaling ratios to different attention heads, with the scaling factor monotonically increasing from "position-aware" heads to "position-unaware" heads. This enables us to improve long-context ability by re-scaling position indices to shorter values while preserving essential knowledge acquired during the pre-training phase. The efficacy of Ms-PoE is substantiated through extensive experiments. By simply re-scaling the indices of positional encoding, Ms-PoE consistently enhances the performance of various LLMs including Llama-2 (Touvron et al., 2023), StableBeluga (Mahan et al.) and Vicuna (Chiang et al., 2023) on the ZeroSCROLLS benchmark (Shaham et al., 2023), achieving a notable average accuracy gain of up to 3.8.

## 2. Background and Related Works

In this section, we provide a concise overview of the background knowledge and recent literature about the generative inference process of Large Language Models (LLMs), their abilities for long-context reasoning, and details of positional encoding.

## 2.1. Generative Inference of LLMs

The generative inference process in LLMs can be categorized into two distinct phases: ① Prefilling Stage: In this initial phase, LLMs receive an input sequence containing detailed instructions that define a specific generation goal. Throughout this stage, intermediate Key and Value embeddings are generated at each layer and stored in memory, commonly referred to as the KV cache. ② Decoding Stage: This phase involves retrieving embeddings from the KV cache to generate new tokens. The decoding process is inherently iterative, where each newly generated token serves as input for the subsequent token generation. In real-world LLM deployment, the cumulative length of input sequences and the subsequently generated text can reach several thousand or even millions of tokens, presenting significant challenges for the LLMs' long-context reasoning capability.

## 2.2. Long Context Reasoning

Two challenges lie ahead for LLMs in handling long-context reasoning tasks. One is to extend the context window to process sentences that exceed the pre-trained window length. Another is the "lost-in-the-window" issue where LLMs will likely overlook the information located in the middle of the sentences.

The reason for the former challenge is that open-source LLMs are usually pre-trained with fixed sequence lengths, such as 4096 for Llama-2 (Touvron et al., 2023). When the sequence length surpasses the predefined context length used in pre-training, LLMs often suffer from performance collapses and thus generate incoherent or fragmented text. Recent efforts to address this issue can be broadly categorized into two streams. Recently, several works have been proposed to address this issue, which can be broadly categorized into two streams. The first one explores from the expansion of positional encoding, with notable contributions including PI (Chen et al., 2023c), CLEX (Chen et al., 2023a), YaRN (Peng et al., 2023), Self-Extend (Jin et al., 2024). On the other hand, some works modify the attention mechanism, such as StreamingLLM (Xiao et al., 2023), LM-Inifinite (Han et al., 2023), H$_2$O (Zhang et al., 2023), TOVA (Oren et al., 2024), Zebra (Song et al., 2023a), and Activation Beacon (Zhang et al., 2024). These approaches have successfully expanded the contextual window with minimal or no additional training overhead.

Despite the extended context window, LLMs still face a significant challenge in long-context inference due to the uneven utilization of lengthy inputs. Liu et al. (2023) conducted a pivotal investigation, revealing that LLMs tend to overlook the middle portion of the input. This bias compromises the practical application of LLMs, as critical information may be located in the middle part of the input, leading to unreliable outputs. To tackle this issue, Peysakhovich & Lerer (2023) introduced 'attention sorting' to reorder inputs, placing critical information at the end. However, this method's reliance on potentially biased attention scores to identify crucial content may compromise its reliability, and the prerequisite knowledge of document count in inputs may affect its effectiveness. Chen et al. (2023d) utilize Attention Buckets, an ensemble approach that combines multiple forward processes with positional modifications. However, this technique necessitates a considerably higher computational cost. Other general approaches for enhancing long-context reasoning include prompt compression (Jiang et al., 2023b), retrieval augmentation (Xu et al., 2023), and inference refinement by constructing memory trees (Chen et al., 2023b) while these approaches typically necessitate extra LLMs' assistance or bring extra computational cost.

## 2.3. Positional Encoding

For effective processing of long contexts, LLMs necessitate the explicit encoding of positional information. Common techniques include absolute positional embedding and relative positional encoding. Absolute positional embedding integrates word embeddings with an additional positional vector based on the token's absolute position, which can be either fixed (Vaswani et al., 2017) or learnable (Devlin et al., 2018; Lan et al., 2019; Clark et al., 2020; Radford et al., 2019; 2018). In contrast, relative positional encoding, increasingly popular in contemporary LLMs, encodes the relative distances between tokens instead of their absolute positions. Notable among these are Rotary Position Embedding (RoPE) (Su et al., 2024) that widely implemented in models like Llama (Touvron et al., 2023), Falcon (Penedo et al., 2023), Mistral (Jiang et al., 2023a), and ALiBi (Press et al., 2021), which used in MPT (Team, 2023).

**RoPE.** The primary goal of RoPE (Su et al., 2024) is to encode positional information such that the inner product of the query and key embeddings inherently contains the relative position information, that is:

$$f(\mathbf{q}_m, m)^T f(\mathbf{k}_n, n) = g(\mathbf{q}_m, \mathbf{k}_n, m - n)$$

Here, $f$ is the positional encoding function applied to the query and key embeddings at positions $m$ and $n$, respectively. To satisfy this condition, the function $f$ is defined as a vector-valued complex function, as follows:

$$f(\mathbf{x}, m) = \mathbf{x}e^{im\theta}$$
$$= [(x_1 + ix_2)e^{im\theta_1}, (x_3 + ix_4)e^{im\theta_2},$$
$$..., (x_{l-1} + ix_l)e^{im\theta_{l/2}}]^T$$

In this equation, $l$ represents the dimension of the embeddings, $\theta_k = 10000^{-2k/l}$, and $i$ is the imaginary unit. For calculating the attention score, RoPE considers the real part of the product, specifically $\text{Re}(f(\mathbf{q}_m, m)^T f(\mathbf{k}_n, n))$.
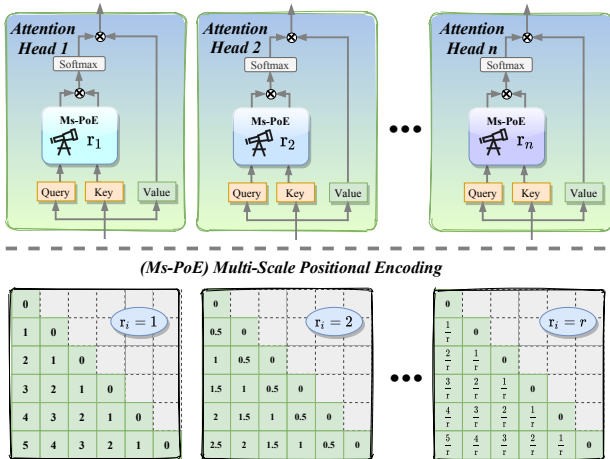
*Figure 2.* Illustration of our Multi-scale Positional Encoding (Ms-PoE) framework. The top figure demonstrates the implementation of Ms-PoE with various scaling ratios in different attention heads, marked with different colors. The bottom figure shows the position details of each head, in which the first matrix ($r_i = 1$) represents the original RoPE.

This approach allows RoPE to effectively integrate relative positional information into the attention mechanism of transformer models.

## 3. Methodology

In this section, we present the details of our Multi-Scale Positional Encoding (Ms-PoE) approach. Section 3.1 demonstrates that the context utilization of LLMs can be directly enhanced by re-scaling the positional information without incurring extra training costs. Then, Section 3.2 analyzes the properties of various attention heads in LLMs. Section 3.3 outlines the detailed pipeline of Ms-PoE.

### 3.1. Positional Re-scaling Improves Context Utilization

Current LLMs tend to neglect information located in the middle of the context, despite its potential relevance. This "lost in the middle" phenomenon likely arises from two contributing factors: (i) Casual Attention, where preceding tokens undergo a higher number of attention processes, leading LLMs to disproportionately favor initial tokens. This phenomenon has been demonstrated in recent research which highlights the pivotal role of the initial tokens in model generation (Han et al., 2023; Xiao et al., 2023), with these starting tokens consistently accumulating higher attention scores (Zhang et al., 2023). (ii) The utilization of RoPE (Su et al., 2024) introduces a long-term decay effect, diminishing the attention score of distantly positioned yet semantically meaningful tokens. The combination of these

factors contributes to LLMs neglecting the context in the middle part. To tackle this issue and improve the context utilization of LLMs, a seemingly unreasonable yet remarkably effective strategy is to down-scale positional information (Song et al., 2023a). Formally, RoPE encodes the position as $f(\mathbf{x}, m) = \mathbf{x}e^{im\theta}$. By substituting the position $m$ with $\frac{m}{r}$, we can force the long-distance tokens to reside in the shorted range, which can potentially alleviate the long-term decay effects by a factor of $r$. In the following sections, we conduct experiments to evaluate how LLMs' context utilization responds to varying re-scaling ratios $r$.

**Details.** Experiments are conducted using Llama-2-7B-Chat (Touvron et al., 2023) and Vicuna-7B (Chiang et al., 2023) on the Multi-Document Question Answering (MDQA) task (Liu et al., 2023). Each question includes ten documents, with only one relevant to the question. By varying the position of the relevant document, we can evaluate LLMs' context utilization properties. For each position of the key document, we calculate the accuracy over 500 samples. And results show in Figure 3 include both the **Average** accuracy over the 10 documents as well as **Gap** accuracy, *i.e.*, the difference between the best and worst accuracy when varying the positions of the relevant document.
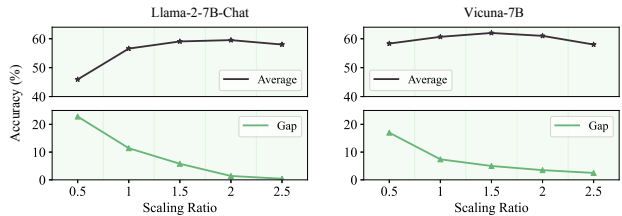


*Figure 3.* Results of the relationship between positional re-scaling and context utilization. The upper curve illustrates the average accuracy when placing the key document in various positions. The bottom curve indicates the gap between the best and worst accuracy.

**Results.** Figure 3 demonstrates that the gap accuracy can be alleviated via appropriate positional re-scaling. Particularly, we see that the Gap between the best and the worst accuracy is greatly reduced when increasing the re-scaling ratio. An enhanced average accuracy can be observed with a scaling ratio equals near 1.5. Additionally, changing the scaling ratio also affects the favored zone of LLMs. With a small scaling ratio (e.g., 0.5), LLMs tend to focus more on the most recent part of the context, while with a large ratio (e.g., 2.5), LLMs favour the beginning part.

**Improving context reasoning via positional re-scaling.** Building upon this, we introduce a plug-and-play treatment for RoPE by re-scaling the position of each token. This approach seamlessly enhances the context utilization of LLMs without requiring additional training or inference

overhead. However, there is a trade-off in terms of LLMs favoring certain context regions. For instance, when $r = 0.5$, LLMs achieve peak accuracy when the relevant document is located at the end of the input, while at the beginning for $r = 1.5$. It remains challenging to decide which re-scaling ratio to use, given that we lack prior knowledge of the location of relevant information in real-world applications. Moreover, as the re-scaling ratio increases, LLMs may face the positional out-of-distribution (O.O.D) issue (Jin et al., 2024; Chen et al., 2023c), where many position values do not directly exist during pretraining (e.g., using $0.1, 0.2, ..., 0.9$ for position when LLMs only recognize $1, 2, ..., 9$ during pretraining), potentially reducing their average reasoning ability. To tackle these challenges, we investigate the head-wise properties of LLMs and propose a multi-scale positional encoding approach.

### 3.2. Position-Aware Head-Wise Re-scaling Ratio

Inspired by recent works that leverage attention patterns to identify most crucial tokens and optimize inference efficiency (Oren et al., 2024; Zhang et al., 2023; Ge et al., 2023), we carry out a preliminary study to investigate the interaction between attention patterns and token positions.

**Details.** We visualize the attention patterns of the most recent query with results collected from Vicuna-7B on the MDQA task, following (Oren et al., 2024). In the same input sample, we manually switch the position of the relevant document from the beginning to the end and illustrate the attention scores across different positions.

**Observation.** We observe the presence of "position-aware" attention heads capable of capturing relevant information even when its position is shifted. As an example, we select the eighth attention head in the fifteenth layer, depicted in the bottom of Figure 4, while consistent observations can be drawn across different layers and input samples. Firstly, most attention scores are near zero and can be ignored, consistent with other studies highlighting high sparsity in attention blocks (Zhang et al., 2023; Likhosherstov et al., 2021; Edelman et al., 2022). For the remaining positions, these "position-aware" attention heads can capture important information across positions, with attention patterns shifting as the position of relevant tokens changes. However, for other attention heads (upper subfigure in Figure 4), they fail to capture relevant tokens and only attend to the beginning and end words, contributing to the "lost-in-the-middle" issue.

Based on this observation, we devise a position-aware strategy to adaptively determine the re-scaling ratio via the inherent properties of different attention heads. For the "position-aware" attention heads, we assign a re-scaling ratio close to one to avoid changing their functionality significantly, as altering them too much could degrade performance due to the positional O.O.D issue. On the other heads, we condense

their position indices to a higher degree, providing more opportunity to alleviate the persistent bias toward the beginning and recent tokens. To identify the properties of $n_h$ attention heads, we introduce a Position-Awareness Score $\mathcal{S}_P \in \mathbf{R}^{n_h}$ formulated as:

$$\mathcal{S}_P = \frac{1}{l} \sum_{i=1}^{l} (A_i \geq \alpha \frac{1}{l} \sum_{i=1}^{l} A_i) \tag{1}$$

In Equation 1, $A$ represents the attention score vector of the most recent query, and $\alpha$ is a hyperparameter determining the threshold of effective attention scores. In all experiments, we default to using $\alpha = 3$, and the corresponding important tokens are highlighted in Figure 4, which are shown in red. In the spirit of numerous studies that investigate the outlier properties in LLMs (Xiao et al., 2023; Lin et al., 2023; Yin et al., 2023), we utilize $\mathcal{S}_P$ to evaluate the ratio of effective attention tokens, where a larger $\mathcal{S}_P$ value implies better positional awareness.

### 3.3. Inference with Multi-Scale Positional Encoding

---
**Algorithm 1** LLM Inference with Ms-PoE
---
**Require:** A pre-trained LLM $f(\theta)$ with positional encoding RoPE, input content $X$, generation length $l$, number of attention heads $n_h$, number of layers $n_l$, re-scaling ratios $\mathbf{r} = [r_1, r_2, \ldots, r_{n_h}]$, calculated in Equation 2.
**Ensure:** Generated text $T$.
 1: Set `Prefilling = True`;
 2: **while** $i < l$ **do**
 3:    **if** `Prefilling` **then**
 4:       **for** k in 1, 2, ..., $n_l$ **do**
 5:          Get the last query $Q_l$ and all key $K$ in layer k;
 6:          $A = \text{Softmax}(Q_{l,*}(K_*)^T)$;
 7:          Calculate Position-Awareness Score $\mathcal{S}_{\mathbf{P}}(A)$ for each head, with Equation 1;
 8:          *# Initial re-scaling ratio in each layer*
 9:          $\mathbf{r}_k = \mathbf{r}[\mathcal{S}_{\mathbf{P}}.\text{argsort}(\text{reverse=True})]$;
10:          Replace RoPE in layer k with Ms-PoE ($\mathbf{r}_k$);
11:       **end for**
12:       `Prefilling = False`;
13:    **end if**
14:    Generate $T$ with implemented Ms-PoE;
15: **end while**
---

The pipeline for utilizing Multi-Scale Positional Encoding (Ms-PoE) in LLM inference is outlined in Algorithm 1. Given a pre-trained LLM, we initially replace the original rotary positional encoding with Ms-PoE. As illustrated in Figure 2, Ms-PoE condenses the positional indices of RoPE and employs different re-scaling ratios for each attention head. The re-scaling ratios are assigned during the prefilling stage, where we first calculate the distribution of attention scores for the most recent query and obtain the
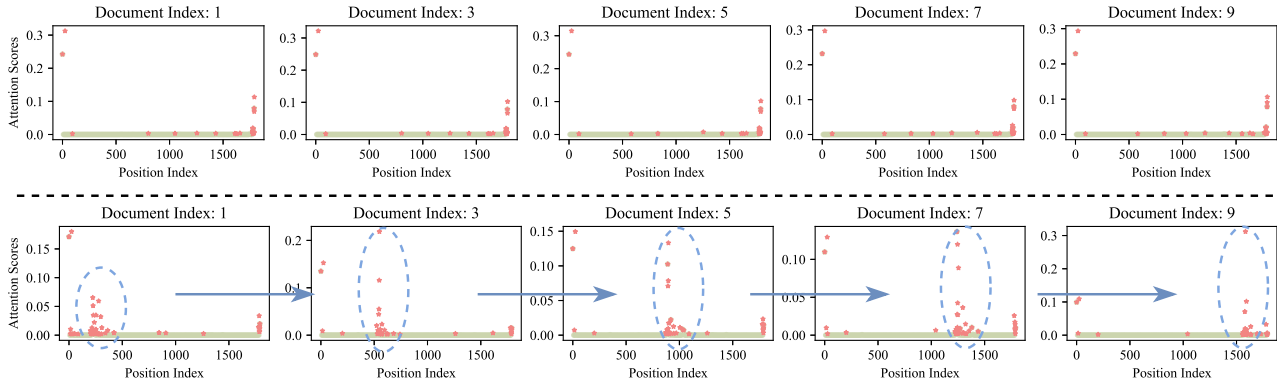
*Figure 4.* Visualization of attention pattern of the most recent query within two different attention heads. **Top:** Results of the 12th attention head in the 15th layer. **Bottom:** Results of the 8th attention head in the 15th layer. The most recent query remains unchanged while varying the position of the crucial document. More examples are reported in Figure 6 in the appendix.

corresponding position-awareness score for each attention head. Larger re-scaling ratios are subsequently allocated to attention heads exhibiting smaller position-awareness scores. And the set of re-scaling ratios **r** defaults to a **linear** range from $1.2$ to $1.8$. For example, the $i$th sorted-head would be using re-scaling ratio

$$r_i = 1.2 + (i-1)(1.8-1.2)/(n_h - 1) \qquad (2)$$

Once the re-scaling ratios are assigned, they remain fixed in the subsequent decoding stage.

## 4. Experiments

The goal of this section is to demonstrate Ms-PoE, a plug-and-play positional encoding capable of enhancing the context utilization of LLMs, and consequently improving the quality of generation across diverse models and downstream reasoning tasks. Our main results can be summarized below.

In Section 4.1, we demonstrate that Ms-PoE consistently enhances reasoning over long contexts for a range of tasks in the ZeroSCROLLS benchmarks (Shaham et al., 2023), all without the need for additional training. Additionally, Ms-PoE exhibits superior performance when compared to other methods in the field, including PI (Chen et al., 2023c) and Self-Extend (Jin et al., 2024).[1] Detailed results of these comparisons are shown in Tables 1 and 2.

In section 4.2, we highlight that Ms-PoE improves the context utilization and achieves consistent improvement when varying the position of critical information within the input context, as shown in Figure 1 & 5.

---

[1] Although the primary intention of PI is to extend sequence lengths, not specifically to tackle the "lost-in-the-middle" problem, we include it in our comparison for its related concept and its approach to positional information handling.

In Section 4.3, we conduct multiple ablation studies to assess the effectiveness of Ms-PoE under different scaling ratios and selection strategies. Results are reported in Table 3 & 4.

### 4.1. Enhanced Generation Quality

We empirically validate the ability of Ms-PoE to enhance long-context reasoning with a noteworthy improvement up to $13.4$ without additional training overhead. Notably, our approach surpasses other competitive baselines, demonstrating improvements from $2.64$ to $43.72$.

**Experimental Setup.** In our experiments, we select seven representative LLMs, including Llama-2-chat-7B and 13B (Touvron et al., 2023), StableBeluga-7B and 13B (Mahan et al.), and Vicuna-7B (Chiang et al., 2023), along with its longer-context version (Vicuna-7B-16K). To comprehensively evaluate the long-context reasoning abilities of LLMs, we choose seven tasks from ZeroSCROLLS (Shaham et al., 2023), spanning all four task categories: ① Document Summarization (Government and SummScreenFD), ② Query-Based Summarization (QMSum and SQuALITY), ③ Question Answering (Qasper and NarrativeQA), and ④ Information Aggregation (BookSumSort). Furthermore, we compare Ms-PoE with other competitive methods on additional generation tasks, including Multi-document Question Answering (MDQA) and Key-Value Retrieval (Liu et al., 2023).

**Main Results.** Table 1 summarizes the main results, yielding several key observations: (i) By simply substituting the original positional encoding module with our Ms-PoE, the performance of LLMs consistently improves across all tasks without additional training, resulting in an average performance enhancement ranging from **0.6** to **3.8**; (ii) These

*Table 1.* Comparsion results on ZeroSCROLLS (Shaham et al., 2023) benchmarks. The evaluation metrics for various tasks are tailored as follows: GovReport, SummScreenFD, QMSum, and SQuALITY utilize the geometric mean of Rouge-1/2/L scores. Qasper and NarrativeQA are assessed through the F1 score, while BookSumSort employs the concordance index.

| Models | Methods | GovReport | SummScreenFD | QMSum | SQuALITY | Qasper | NarrativeQA | BookSumSort | Average |
|---|---|---|---|---|---|---|---|---|---|
| Llama-2-7B-Chat | Baseline | 16.8 | 14.1 | 15.2 | 19.5 | 21.9 | 14.4 | 3.1 | 15.0 |
| Llama-2-7B-Chat | Ours | **17.7 (+0.9)** | **14.2 (+0.1)** | **15.8 (+0.6)** | **19.9 (+0.4)** | **25.1 (+3.2)** | **17.7 (+3.3)** | **5.8 (+2.7)** | **16.6 (+1.6)** |
| Llama-2-13B-Chat | Baseline | 15.4 | 12.3 | 15.1 | 18.9 | 19.0 | 15.0 | 5.7 | 14.5 |
| Llama-2-13B-Chat | Ours | **16.5 (+1.1)** | **13.1 (+0.8)** | **15.5 (+0.4)** | **19.2 (+0.3)** | **20.8 (+1.8)** | **17.0 (+2.0)** | **5.9 (+0.2)** | **15.4 (+0.9)** |
| StableBeluga-7B | Baseline | 14.9 | 13.8 | 14.7 | 17.9 | 28.1 | 16.8 | 9.2 | 16.5 |
| StableBeluga-7B | Ours | **16.6 (+1.7)** | **14.2 (+0.4)** | **15.2 (+0.5)** | **18.7 (+0.8)** | **36.9 (+8.8)** | **18.0 (+1.2)** | **14.2 (+5.0)** | **19.1 (+2.6)** |
| StableBeluga-13B | Baseline | 5.7 | 7.1 | 12.9 | 13.3 | 19.2 | 13.4 | 4.8 | 10.9 |
| StableBeluga-13B | Ours | **7.4 (+1.7)** | **7.4 (+0.3)** | 12.8 (-0.1) | 13.2 (-0.1) | **20.8 (+1.6)** | 13.4 (+0) | **5.6 (+0.8)** | **11.5 (+0.6)** |
| Vicuna-7B | Baseline | 16.2 | 13.7 | 15.1 | 18.9 | 24.3 | 13.7 | 3.3 | 15.0 |
| Vicuna-7B | Ours | **20.2 (+4.0)** | **14.5 (+1.8)** | **15.4 (+0.3)** | **19.8 (+0.9)** | **34.7 (+13.4)** | **16.2 (+2.5)** | **10.5 (+7.2)** | **18.8 (+3.8)** |
| Vicuna-7B-16K | Baseline | 20.2 | 13.9 | 16.2 | 20.1 | 32.3 | 18.8 | 29.9 | 21.6 |
| Vicuna-7B-16K | Ours | **21.4 (+1.2)** | **14.3 (+0.4)** | **16.2 (+0)** | **20.2 (+0.1)** | **37.8 (+5.5)** | **21.0 (+2.2)** | **43.3 (+13.4)** | **24.9 (+3.3)** |

improvements hold consistently across different model sizes of 7 billion and 13 billion parameters; (iii) The efficacy extends to LLMs with varying sequence lengths, such as Vicuna-7B and its extended version, Vicuna-7B-16K, both showing improvements from **3.3** to **3.8**.

*Table 2.* Comparsion results with other competitive methods on MDQA and Key-Value Retrieval. Results are reported in accuracy.

| Models | Methods | MDQA | | | | | |
|---|---|---|---|---|---|---|---|
| | | 1 | 3 | 5 | 7 | 10 | Average |
| | Baseline | 64.0 | 61.0 | 57.4 | 58.4 | 64.8 | 61.12 |
| | PI | 65.2 | 62.4 | 60.0 | 60.4 | 64.0 | 62.40 |
| Vicuna-7B | Self-Extend | 64.7 | 63.7 | 61.4 | 59.8 | 62.0 | 62.32 |
| | Ms-PoE | **65.6** | **64.2** | **63.0** | **65.2** | **67.2** | **65.04** |

| Models | Methods | Key-Value Retrival | | | | | |
|---|---|---|---|---|---|---|---|
| | | 1 | 15 | 30 | 40 | 50 | Average |
| | Baseline | 92.0 | 25.8 | 8.0 | 25.4 | 30.0 | 36.24 |
| | PI | 96.4 | 76.4 | 61.4 | 64.6 | **57.8** | 67.60 |
| Vicuna-7B | Self-Extend | 88.6 | 63.8 | 76.2 | 59.4 | 42.0 | 66.00 |
| | Ms-PoE | **97.0** | **83.4** | **75.0** | **86.6** | 57.8 | **79.96** |

**Outperform other competitive methods.** We conduct a thorough comparison between Ms-PoE and other competitive methods, including Positional Interpolation (PI) (Chen et al., 2023c) and Self-Extend (Jin et al., 2024), both of which modify position indices without utilizing head-wise properties. For PI, we employ the scaling ratio as the average value of our method while for Self-Extend, we set the group size as 2 with the local window size as 1024. The results presented in Table 2 consistently showcase the superiority of our approach over other baselines, demonstrating improvements of up to 3.92 and 43.72 for MDQA and Key-Value Retrieval, respectively. Such improvements might come from two primary factors. Firstly, the incorporation of head-wise properties offers a more adaptive strategy for positional modification. Secondly, our approach enhances the general context utilization ability. Notably, our approach demonstrates superiority even when the core document or key is positioned at the end of the input, surpassing other
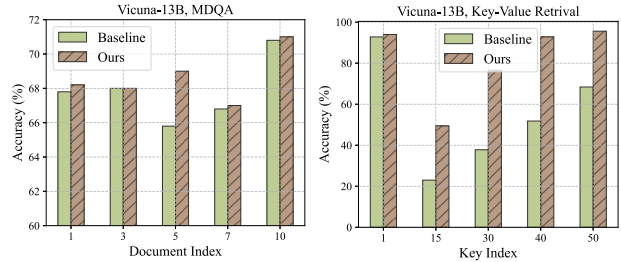


*Figure 5.* Comparison results for the multi-document question answering (MDQA) and key-value retrieval (KV retrieval) tasks. Each subfigure depicts the comparison when varying the position of critical information from the beginning to the end. For Vicuna-7B, please refer to Figure 1.

baselines with improvements ranging from 2.4 to 27.8. This performance surpasses the recent work (Peysakhovich & Lerer, 2023), which addresses the "lost-in-the-middle" effect by reordering key documents and placing them at the end of the input. When the identified core document is already located at the recent area, such method can not gain further improvements, while our approach offers a *fine-grained* strategy to improve context utilization.

**4.2. Superior Context Utilization**

We assess the context utlization ability of our approaches on two tasks, including multi-document question answering (MDQA) and key-value retrieval (KV retrieval) tasks from (Liu et al., 2023). Such tasks provide a good input structure and offers the flexibility to switch the position of crusial information, thus evaluate the context utilization ability of LLMs.

**Experimental Setup.** In the MDQA task, each input sample comprises ten documents and one question, with only

one document being relevant to the question. For the KV retrieval tasks, there are 50 key-value pairs with one question querying the value of the chosen key. In both tasks, we systematically switch the important document or key-value pair from the beginning to the end and report the accuracy of the generated context. All results are averaged across 500 samples. The **Gap** accuracy metric is employed to assess the context utilization ability of LLMs, defined as the gap between the best and worst accuracy when varying the position of important information.

**Main Results.** As depicted in Figure 5 and 1, Ms-PoE demonstrates consistent improvement across different models, tasks and critical positions. Even when the important information exists in the sweet region (beginning and end) of the input, Ms-PoE achieves significant performance improvements ranging from 3% to 6%, highlighting its efficacy in enhancing generation quality. Moreover, the "lost-in-the-middle" issue is notably alleviated, with Ms-PoE quantitatively reducing the gap accuracy by approximately 2% to 4%, showcasing improved context utilization.

### 4.3. Ablation Study and More Investigation

This section conducts a further evaluation of the effectiveness of Ms-PoE by addressing the following questions: *Q1:* How does the effectiveness of Ms-PoE relate to the head-wise selection strategy of the scaling ratio? *Q2:* How does the model perform with different scaling ratios?

*A1:* **Positional awareness metrics achieve superior performance compared to other strategies.** For a set of scaling ratios $\mathbf{r} \in R^{n_h}$, where $n_h$ is the number of attention heads, and using scaling ratios linearly ranging from 1.2 to 1.8, we evaluate various strategies for assigning these ratios to different attention heads. These strategies include: ① `Random`, which randomly assigns the scaling ratios to each head within each layer; ② `Sequential`, performing the assignment based on the original head order; ③ `Entropy`, where we follow metrics measuring the sparsity level of attention scores (Tian et al., 2023). Larger entropy implies less sparse attention scores, indicating the model attends to more tokens rather than just the beginning and end words, so we assign a scaling ratio near to 1, and vice versa for larger ratios. Results in Table 3 demonstrate that the proposed position-awareness effectively captures the head-wise properties of LLMs, enhancing performance when critical information is located at various positions—beginning, middle, or end. This leads to an average accuracy gain of 3.2 (65.3 v.s. 62.1).

*A2:* **Ablation study of the scaling ratios.** We first examined the effect of uniform scaling ratios across all heads on model performance. Our findings, outlined in Table 4, indicate that adjusting the scaling ratio between 0.5 and 2.5 can

*Table 3.* Ablation results of different ordering metrics. Experiments are conducted on Multi-Documents Question Answering task with the Vicuna-7B model.

| Methods | Begin | Middle | End | Average |
|---|---|---|---|---|
| Baseline | 64.0 | 57.4 | 64.8 | 62.1 |
| Random | 64.5 | 55.0 | 65.5 | 61.7 |
| Sequential | 60.5 | 54.5 | 58.5 | 57.8 |
| Entropy | 63.5 | 59.5 | 64.0 | 62.3 |
| Position-Awareness | **65.6** | **63.0** | **67.2** | **65.3** |

*Table 4.* Ablation results of the condensing ratios. Experiments are conducted on Multi-Documents Question Answering task with the Vicuna-7B model.

| Scaling Ratio | Begin | Middle | End | Average |
|---|---|---|---|---|
| 1 | 64.0 | 57.4 | 64.8 | 62.1 |
| 0.5 | 56.0 | 51.0 | 68.0 | 58.3 |
| 1.5 | 65.2 | 60.0 | 64.0 | 63.1 |
| 2 | 61.5 | 59.0 | 62.5 | 61.0 |
| 2.5 | 59.5 | 57.5 | 57.0 | 58.0 |
| $0.8 \rightarrow 2.2$ | 53.5 | 59.5 | 67.5 | 60.2 |
| $1 \rightarrow 2$ | 61.0 | 57.0 | 63.0 | 60.3 |
| $1.2 \rightarrow 1.8$ | **65.6** | **63.0** | **67.2** | **65.3** |
| $1.4 \rightarrow 1.6$ | 65.5 | 59.0 | 63.0 | 62.5 |

significantly enhance generative performance and mitigate the "lost-in-the-middle" effect by 1.0% (63.1% *v.s.* 62.1%), particularly with a ratio of 1.5. Further testing with an average ratio of 1.5 across all heads revealed that an optimal range exists between 1.2 and 1.8, leading to an additional 2.2% (65.3% *v.s.* 63.1%) accuracy improvement with our approach, Ms-PoE. Based on these results, we established these ratios as our experimental standard.

## 5. Conclusion

In this paper, we present a plug-and-play strategy designed to address the "lost-in-the-middle" challenge observed in LLMs. This challenge stems from the persistent bias exhibited by LLMs towards the beginning and local content within the input, leading to the neglect of crucial information in the middle. Our investigation reveals the effects of position indice rescaling and the head-wise position-awareness property, leading to the introduction of Multi-scale Positional Encoding (Ms-PoE). This approach enhances the capability of LLMs to effectively capture information in the middle of the context without the need for additional fine-tuning. Comprehensive experiments conducted on Zero-SCROLLS benchmarks, multi-document question-answering tasks, and key-value retrieval tasks con-

firm the effectiveness of Ms-PoE.

## 6. Impact Statement

The introduction of Multi-scale Positional Encoding (Ms-PoE) offers a simple yet effective approach to tackling the challenge of processing long contexts in large language models (LLMs). This enhancement significantly improves LLMs' capabilities to understand and reason over extensive textual contexts. Ms-PoE can potentially impact numerous fields where processing long-context data is crucial, such as in analyzing vast amounts of case law or detailed patient histories. However, as we advance the boundaries of Machine Learning and Artificial Intelligence, making them increasingly powerful and applicable, it comes with potential challenges and problems that must be carefully addressed. For instance, there is a risk of exacerbating existing biases within AI, as processing longer contexts could amplify the biases present in the training data. Additionally, the misuse of predictive insights derived from LLMs presents a significant worry. To mitigate these concerns, we advocate for ongoing ethical evaluations and the development of guidelines to ensure that the applications of LLM advancements contribute positively to society.

## 7. Acknowledgements

We thank Dr. Yuandong Tian for interesting discussions on this work.

## References

Chen, G., Li, X., Meng, Z., Liang, S., and Bing, L. Clex: Continuous length extrapolation for large language models. *arXiv preprint arXiv:2310.16450*, 2023a.

Chen, H., Pasunuru, R., Weston, J., and Celikyilmaz, A. Walking down the memory maze: Beyond context limit through interactive reading. *arXiv preprint arXiv:2310.05029*, 2023b.

Chen, S., Wong, S., Chen, L., and Tian, Y. Extending context window of large language models via positional interpolation. *arXiv preprint arXiv:2306.15595*, 2023c.

Chen, Y., Lv, A., Lin, T.-E., Chen, C., Wu, Y., Huang, F., Li, Y., and Yan, R. Fortify the shortest stave in attention: Enhancing context awareness of large language models for effective tool use. *arXiv preprint arXiv:2312.04455*, 2023d.

Chiang, W.-L., Li, Z., Lin, Z., Sheng, Y., Wu, Z., Zhang, H., Zheng, L., Zhuang, S., Zhuang, Y., Gonzalez, J. E., Stoica, I., and Xing, E. P. Vicuna: An open-source chatbot impressing gpt-4 with 90%* chatgpt quality, March 2023. URL https://lmsys.org/blog/2023-03-30-vicuna/.

Clark, K., Luong, M.-T., Le, Q. V., and Manning, C. D. Electra: Pre-training text encoders as discriminators rather than generators. *arXiv preprint arXiv:2003.10555*, 2020.

Dao, T. FlashAttention-2: Faster attention with better parallelism and work partitioning. 2023.

Devlin, J., Chang, M.-W., Lee, K., and Toutanova, K. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*, 2018.

Du, X., Liu, M., Wang, K., Wang, H., Liu, J., Chen, Y., Feng, J., Sha, C., Peng, X., and Lou, Y. Classeval: A manually-crafted benchmark for evaluating llms on class-level code generation. *arXiv preprint arXiv:2308.01861*, 2023.

Edelman, B. L., Goel, S., Kakade, S., and Zhang, C. Inductive biases and variable creation in self-attention mechanisms. In *International Conference on Machine Learning*, pp. 5793–5831. PMLR, 2022.

Ge, S., Zhang, Y., Liu, L., Zhang, M., Han, J., and Gao, J. Model tells you what to discard: Adaptive kv cache compression for llms. *arXiv preprint arXiv:2310.01801*, 2023.

Guo, M.-H., Lu, C.-Z., Hou, Q., Liu, Z., Cheng, M.-M., and Hu, S.-M. Segnext: Rethinking convolutional attention design for semantic segmentation. *Advances in Neural Information Processing Systems*, 35:1140–1156, 2022.

Han, C., Wang, Q., Xiong, W., Chen, Y., Ji, H., and Wang, S. Lm-infinite: Simple on-the-fly length generalization for large language models. *arXiv preprint arXiv:2308.16137*, 2023.

Ivgi, M., Shaham, U., and Berant, J. Efficient long-text understanding with short-text models. *Transactions of the Association for Computational Linguistics*, 11:284–299, 2023.

Jacobs, S. A., Tanaka, M., Zhang, C., Zhang, M., Song, S. L., Rajbhandari, S., and He, Y. Deepspeed ulysses: System optimizations for enabling training of extreme long sequence transformer models. *arXiv preprint arXiv:2309.14509*, 2023.

Jiang, A. Q., Sablayrolles, A., Mensch, A., Bamford, C., Chaplot, D. S., Casas, D. d. l., Bressand, F., Lengyel, G., Lample, G., Saulnier, L., et al. Mistral 7b. *arXiv preprint arXiv:2310.06825*, 2023a.

Jiang, A. Q., Sablayrolles, A., Roux, A., Mensch, A., Savary, B., Bamford, C., Chaplot, D. S., Casas, D. d. l., Hanna, E. B., Bressand, F., et al. Mixtral of experts. *arXiv preprint arXiv:2401.04088*, 2024.

Jiang, H., Wu, Q., Luo, X., Li, D., Lin, C.-Y., Yang, Y., and Qiu, L. Longllmlingua: Accelerating and enhancing llms in long context scenarios via prompt compression. *arXiv preprint arXiv:2310.06839*, 2023b.

Jin, H., Han, X., Yang, J., Jiang, Z., Liu, Z., Chang, C.-Y., Chen, H., and Hu, X. Llm maybe longlm: Self-extend llm context window without tuning. *arXiv preprint arXiv:2401.01325*, 2024.

Junqing, H., Kunhao, P., Xiaoqun, D., Zhuoyang, S., Yibo, L., Yuxin, L., Hao, W., Qianguo, S., Songxin, Z., Zejian, X., et al. Never lost in the middle: Improving large language models via attention strengthening question answering. *arXiv preprint arXiv:2311.09198*, 2023.

Kryściński, W., Rajani, N., Agarwal, D., Xiong, C., and Radev, D. Booksum: A collection of datasets for long-form narrative summarization. *arXiv preprint arXiv:2105.08209*, 2021.

Lan, Z., Chen, M., Goodman, S., Gimpel, K., Sharma, P., and Soricut, R. Albert: A lite bert for self-supervised learning of language representations. *arXiv preprint arXiv:1909.11942*, 2019.

Li, J., Wang, M., Zheng, Z., and Zhang, M. Loogle: Can long-context language models understand long contexts? *arXiv preprint arXiv:2311.04939*, 2023.

Likhosherstov, V., Choromanski, K., and Weller, A. On the expressive power of self-attention matrices. *arXiv preprint arXiv:2106.03764*, 2021.

Lin, J., Tang, J., Tang, H., Yang, S., Dang, X., and Han, S. Awq: Activation-aware weight quantization for llm compression and acceleration. *arXiv preprint arXiv:2306.00978*, 2023.

Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., and Liang, P. Lost in the middle: How language models use long contexts. *arXiv preprint arXiv:2307.03172*, 2023.

Mahan, D., Carlow, R., Castricato, L., Cooper, N., and Laforte, C. Stable beluga models. URL [https://huggingface.co/stabilityai/StableBeluga2](https://huggingface.co/stabilityai/StableBeluga2).

Oren, M., Hassid, M., Adi, Y., and Schwartz, R. Transformers are multi-state rnns. *arXiv preprint arXiv:2401.06104*, 2024.

Penedo, G., Malartic, Q., Hesslow, D., Cojocaru, R., Cappelli, A., Alobeidli, H., Pannier, B., Almazrouei, E., and Launay, J. The RefinedWeb dataset for Falcon LLM: outperforming curated corpora with web data, and web data only. *arXiv preprint arXiv:2306.01116*, 2023. URL https://arxiv.org/abs/2306.01116.

Peng, B., Quesnelle, J., Fan, H., and Shippole, E. Yarn: Efficient context window extension of large language models. *arXiv preprint arXiv:2309.00071*, 2023.

Peysakhovich, A. and Lerer, A. Attention sorting combats recency bias in long context language models. *arXiv preprint arXiv:2310.01427*, 2023.

Press, O., Smith, N. A., and Lewis, M. Train short, test long: Attention with linear biases enables input length extrapolation. *arXiv preprint arXiv:2108.12409*, 2021.

Radford, A., Narasimhan, K., Salimans, T., Sutskever, I., et al. Improving language understanding by generative pre-training. 2018.

Radford, A., Wu, J., Child, R., Luan, D., Amodei, D., Sutskever, I., et al. Language models are unsupervised multitask learners. *OpenAI blog*, 1(8):9, 2019.

Ré, C., Dao, T., Fu, D., and Goel, K. Can longer sequences help take the next leap in ai?, June 2022. URL https://hazyresearch.stanford.edu/blog/2022-06-09-longer-sequences-next-leap-ai. Accessed: 2024-01-29.

Shaham, U., Ivgi, M., Efrat, A., Berant, J., and Levy, O. Zeroscrolls: A zero-shot benchmark for long text understanding. *arXiv preprint arXiv:2305.14196*, 2023.

Song, K., Wang, X., Cho, S., Pan, X., and Yu, D. Zebra: Extending context window with layerwise grouped local-global attention. *arXiv preprint arXiv:2312.08618*, 2023a.

Song, S. L., Kruft, B., Zhang, M., Li, C., Chen, S., Zhang, C., Tanaka, M., Wu, X., Rasley, J., Awan, A. A., et al. Deepspeed4science initiative: Enabling large-scale scientific discovery through sophisticated ai system technologies. *arXiv preprint arXiv:2310.04610*, 2023b.

Su, J., Ahmed, M., Lu, Y., Pan, S., Bo, W., and Liu, Y. Roformer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568:127063, 2024.

Szegedy, C., Liu, W., Jia, Y., Sermanet, P., Reed, S., Anguelov, D., Erhan, D., Vanhoucke, V., and Rabinovich, A. Going deeper with convolutions. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 1–9, 2015.

Szegedy, C., Vanhoucke, V., Ioffe, S., Shlens, J., and Wojna, Z. Rethinking the inception architecture for computer vision. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 2818–2826, 2016.

Tay, Y., Dehghani, M., Abnar, S., Shen, Y., Bahri, D., Pham, P., Rao, J., Yang, L., Ruder, S., and Metzler, D. Long range arena: A benchmark for efficient transformers. In *International Conference on Learning Representations*, 2020.

Team, M. N. Introducing mpt-7b: A new standard for open-source, commercially usable llms, 2023. URL www.mosaicml.com/blog/mpt-7b. Accessed: 2023-05-05.

Tian, Y., Wang, Y., Zhang, Z., Chen, B., and Du, S. Joma: Demystifying multilayer transformers via joint dynamics of mlp and attention. *arXiv preprint arXiv:2310.00535*, 2023.

Touvron, H., Martin, L., Stone, K., Albert, P., Almahairi, A., Babaei, Y., Bashlykov, N., Batra, S., Bhargava, P., Bhosale, S., et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.

Varadi, M., Anyango, S., Deshpande, M., Nair, S., Natassia, C., Yordanova, G., Yuan, D., Stroe, O., Wood, G., Laydon, A., et al. Alphafold protein structure database: massively expanding the structural coverage of protein-sequence space with high-accuracy models. *Nucleic acids research*, 50(D1):D439–D444, 2022.

Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., and Polosukhin, I. Attention is all you need. *Advances in neural information processing systems*, 30, 2017.

Wang, W., Dong, L., Cheng, H., Liu, X., Yan, X., Gao, J., and Wei, F. Augmenting language models with long-term memory. *arXiv preprint arXiv:2306.07174*, 2023.

Xiao, G., Tian, Y., Chen, B., Han, S., and Lewis, M. Efficient streaming language models with attention sinks. *arXiv preprint arXiv:2309.17453*, 2023.

Xu, P., Ping, W., Wu, X., McAfee, L., Zhu, C., Liu, Z., Subramanian, S., Bakhturina, E., Shoeybi, M., and Catanzaro, B. Retrieval meets long context large language models. *arXiv preprint arXiv:2310.03025*, 2023.

Yin, L., Wu, Y., Zhang, Z., Hsieh, C.-Y., Wang, Y., Jia, Y., Pechenizkiy, M., Liang, Y., Wang, Z., and Liu, S. Outlier weighed layerwise sparsity (owl): A missing secret sauce for pruning llms to high sparsity, 2023.

Zhang, P., Liu, Z., Xiao, S., Shao, N., Ye, Q., and Dou, Z. Soaring from 4k to 400k: Extending llm's context with activation beacon. *arXiv preprint arXiv:2401.03462*, 2024.

Zhang, Y., Ni, A., Mao, Z., Wu, C. H., Zhu, C., Deb, B., Awadallah, A. H., Radev, D., and Zhang, R. Summ^n: A multi-stage summarization framework for long input dialogues and documents. *arXiv preprint arXiv:2110.10150*, 2021.

Zhang, Z., Sheng, Y., Zhou, T., Chen, T., Zheng, L., Cai, R., Song, Z., Tian, Y., Ré, C., Barrett, C., et al. H _2 o: Heavy-hitter oracle for efficient generative inference of large language models. *arXiv preprint arXiv:2306.14048*, 2023.

Zheng, Q., Xia, X., Zou, X., Dong, Y., Wang, S., Xue, Y., Wang, Z., Shen, L., Wang, A., Li, Y., Su, T., Yang, Z., and Tang, J. Codegeex: A pre-trained model for code generation with multilingual evaluations on humaneval-x. In *KDD*, 2023.

Zhong, M., Liu, Y., Xu, Y., Zhu, C., and Zeng, M. Dialoglm: Pre-trained model for long dialogue understanding and summarization. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 36, pp. 11765–11773, 2022.

# A. More Experiment Results

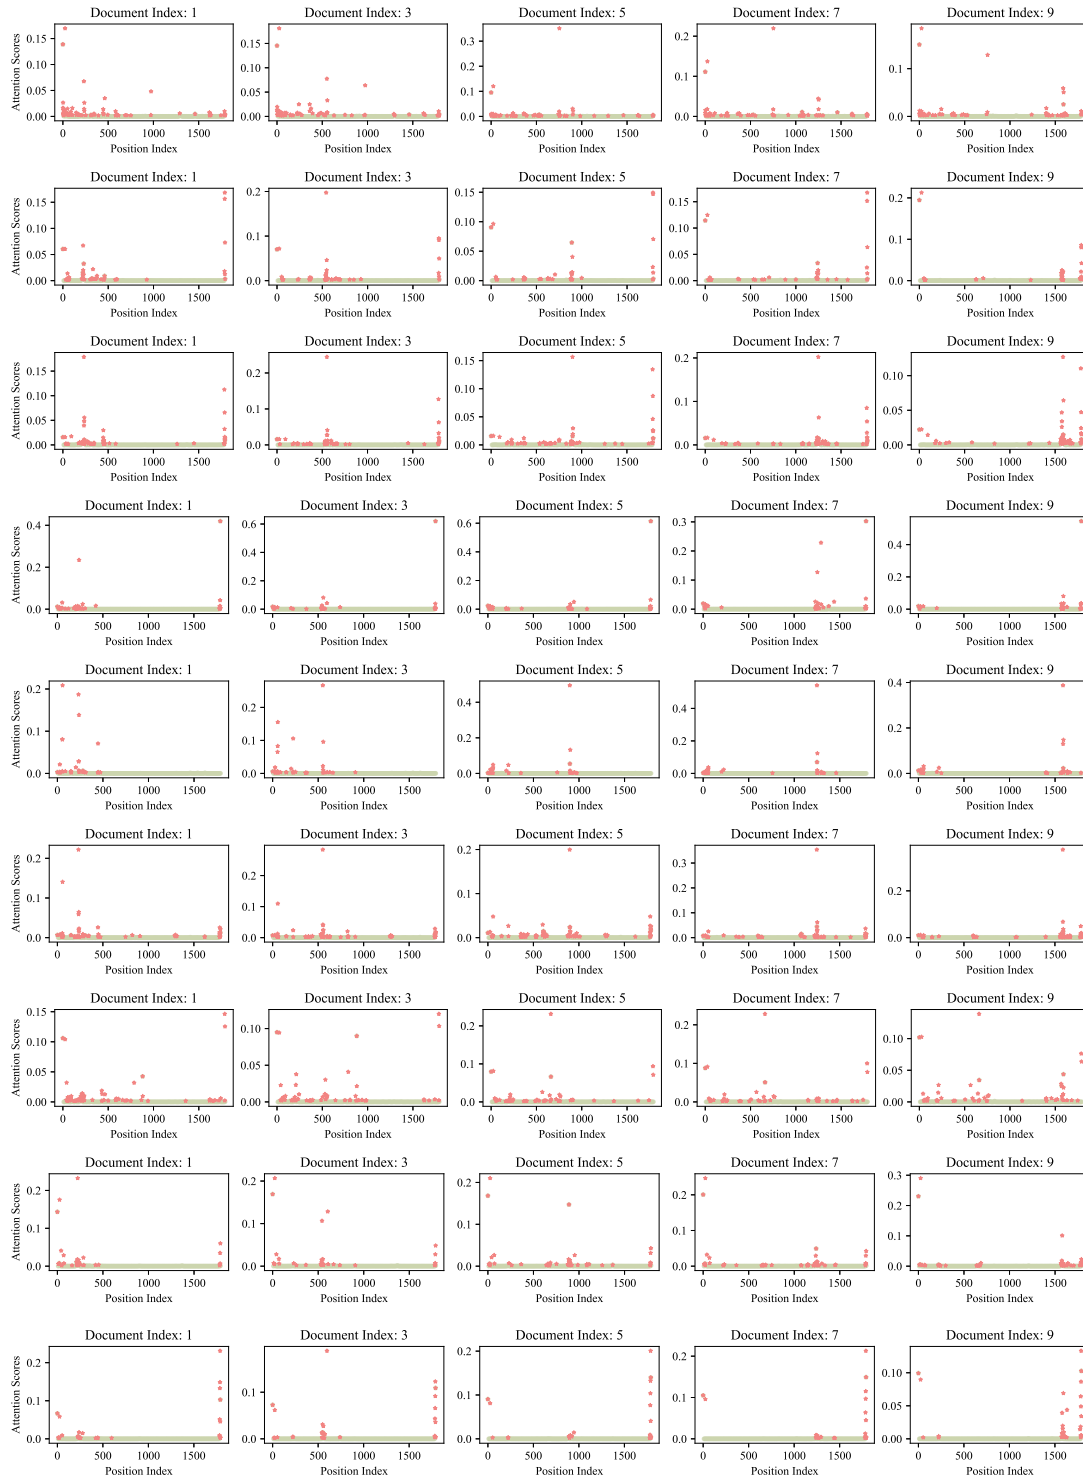## A.1. Position-Aware Attention Heads



*Figure 6.* Visualization of "position-aware" attention heads. Each row contains the attention pattern for the same heads when varying the key documents within the inputs.

Figure 6 illustrates the attention patterns of "position-aware" heads. Each row represents the attention pattern of the same head. As the key document is positioned from the beginning to the end, the attention peak gradually shifts, indicating robust positional awareness. It's important to note that we randomly selected 9 attention heads with these "position-aware" properties, and these results were validated with different input samples and layers.